

# Sistemas Operativos Avanzados

---

[vpuente@unican.es](mailto:vpuente@unican.es)

[www.ce.unican.es](http://www.ce.unican.es)

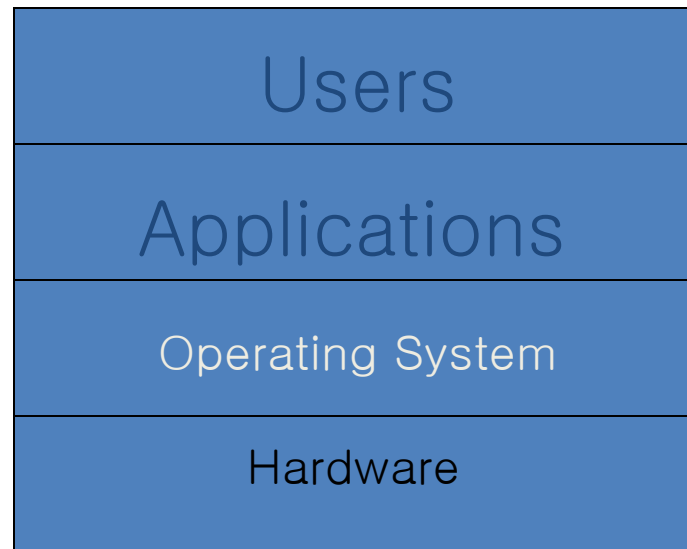


# What is an Operating System?

Operating System (OS):

Software that converts hardware into a useful form for applications

Not easy to define precisely...



# What DOES OS Provide?

- ▣ Role #1: Abstraction - Provide standard library for resources
- ▣ What is a resource?
  - ◆ Anything valuable (e.g., CPU, memory, disk, I/O device)
- ▣ What abstraction does modern OS typically provide for each resource?
  - ◆ CPU:
    - process and/or thread
  - ◆ Memory:
    - address space
  - ◆ Disk:
    - files
- ▣ Advantages of OS providing abstraction?
  - ◆ Allow applications to reuse common facilities
  - ◆ Make different devices look the same
  - ◆ Provide higher-level or more useful functionality
- ▣ Challenges
  - ◆ What are the correct abstractions?
  - ◆ How much of hardware should be exposed?

# What DOES OS Provide?

- ▣ Role #2: Resource management – Share resources well
- ▣ Advantages of OS providing resource management?
  - ◆ Protect applications from one another
  - ◆ Provide efficient access to resources (cost, time, energy)
  - ◆ Provide fair access to resources
- ▣ Challenges
  - ◆ What are the correct mechanisms?
  - ◆ What are the correct policies?

# OS Organization

- ▣ How to cover all the topics relevant to operating systems?

# Three PIECES

- ▣ Virtualization:
  - ◆ Make each application believe it has each resource to itself
  
- ▣ Concurrency:
  - ◆ Events are occurring simultaneously and may interact with one another
  
- ▣ Persistence: Access information permanently
  - ◆ Lifetime of information is longer than lifetime of any one process
  - ◆ Machine may be rebooted, machine may lose power or crash unexpectedly

# Advanced Topics (beyond our reach)

- ▣ Current systems
  - ◆ Multiprocessors
  - ◆ Networked and distributed systems
  - ◆ Virtual machines
  - ◆ Containers
  - ◆ ...
  
- ▣ Many of the pushed by the explosive demand (a.k.a. Massive complexity under constrained cost)
  
- ▣ This is the support of the world: it will keep changing ...
  
- ▣ Some of them covered in SVS (M1679)

# Why study Operating Systems?

- ▣ Build, modify, or administer an operating system
- ▣ Understand system performance
  - ◆ Behavior of OS impacts entire machine
  - ◆ Tune workload performance
  - ◆ Apply knowledge across many layers
    - Computer architecture, programming languages, data structures and algorithms, and performance modeling
- ▣ Fun and challenging to understand large, complex systems
- ▣ **Is the glue that “holds” all the ideas in place**



# Approach

- We will follow “Operating System: Three Easy Pieces” (OSTEP) style
  - ◆ From the **bottom** concepts to state-of-the-art approaches
  - ◆ Eminently **practical** style: all supported by “simulators” and simple coding examples
  - ◆ Assumes some **basic** knowledge in architecture, C, assembler and system administration
  - ◆ More than just a textbook...
  
- Structure
  - ◆ The three parts are split in small *pieces* (~40 in the book)
  - ◆ Each chapter is built over the previous one (can't miss the beat)
  - ◆ Each chapter has attached a “Homework” to reinforce the : from using python simulators to write small pieces of code ( C )
  - ◆ 5 + 1 Labs, developed on top of **xv6**

# Lecture/Lab structure

- ▣ We mix dynamically both
  - ◆ The real thing is that there is no separation between “theory” and “lab”
  
- ▣ Sessions of:
  - ◆ 1<sup>st</sup> hour: Introduction to the topic
  - ◆ 2<sup>nd</sup> hour: Introduce/develop of Labs
  - ◆ Personal work (out the lab): 6 hours (labs and homework)
  - ◆ 10 hours/week
  - ◆ Strict schedule
  
- ▣ Although the original course/book is designed for 15-week semester (150h work), we will need to drop some details or advanced topics (and half of the labs)

# Material

- ▣ Available in <http://www.ce.unican.es/>
- ▣ All written material will be in “English”
  - ◆ Lecture notes, Homework/Lab guides, etc....
- ▣ *Git* as communication “device”: all material will be delivered via <http://gitlab.com> (lab work) and <http://github.com> (lecture notes)
  - ◆ An e-mail inviting to join the course project will be sent to unican account
  - ◆ Slides, labs, other reference material is there
  - ◆ It uses “git” to have a “time-track”
    - Lecture notes updates
    - Additional material
- ▣ Use git to allow you and me “track” your personal work
- ▣ <http://piazza.com> is a great tool to resolve issues and collaborate (support **anonymous** questions!)

# Book (ostep.org)

This book is **and will always be free** in PDF form, as seen below. For those of you wishing to **BUY** a copy, please consider the following:



- **Lulu Hardcover (v1.00)**: this may be the best printed form of the book (it really looks pretty good), but it is also the most expensive way to obtain *the black book* of operating systems (a.k.a. *the comet book* or *the asteroid book* according to students). Now just: **\$38.00**
- **Lulu Softcover (v1.00)**: this way is pretty great too, if you like to read printed material but want to save a few bucks. Now just: **\$22.00**
- **Amazon Softcover (v1.00)**: Same book as softcover above, but printed through Amazon CreateSpace. Now just: **\$25.90** (but works with Prime shipping)
- **Downloadable PDF (v1.00)**: this is a nice convenience and adds things like a hyperlinked table of contents, index of terms, lists of hints, tips, systems advice, and a few other things not seen in the free version, all in one massive DRM-free PDF. Once purchased, you will always be able to get the latest version. Just: **\$10.00**
- **Kindle**: Really, just the PDF and does not include all the bells and whistles common in e-pub books.

**New Partnership:** We have a new partnership with [Educative](#); they offer a way to take an OS course (based on OSTEP) through their platform at a low cost. However, don't worry: the book on this website is and will always remain free.

**Lulu Discount Codes:** These always exist in some form, look around for one?

**Warning:** Some resellers on Amazon buy old versions of the books and claim to sell them as "new" on Amazon (click [here](#) for an example); buy from them at your own risk. In general, buy either directly from Lulu.com or Amazon.com (not a reseller). For Amazon, go to [this page](#) and look for Seller Information to be Amazon.com.

Can't bear to go out in public without OSTEP? How about an [Operating Systems: Three Easy Pieces T-shirt](#) or [laptop sticker](#) or [bathmat](#) or [blanket](#) or [mug](#) or [check out the whole store](#)?

**Donate:** By popular demand, another way to support this site and its contents: **donate!** Click to donate [\\$1](#) - [\\$10](#) - [\\$20](#) - [\\$50](#) - or click [here](#) to donate any amount you want! Your donation helps keep this book going. Think about it: if everyone who came to this website donated just one dollar, we'd have at least three dollars. Thanks!

Another way to help the book out: cite it! Here is the [BiBTeX entry \(seen below\)](#); you can also link to the site of the [best free operating systems book](#) on the market.

## Operating Systems: Three Easy Pieces

Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau  
Arpaci-Dusseau Books  
August, 2018 (Version 1.00)

And now, the free online form of the book, in chapter-by-chapter form (now with chapter numbers!):

Intro	Virtualization		Concurrency	Persistence	Security
Preface	3 <i>Dialogue</i>	12 <i>Dialogue</i>	25 <i>Dialogue</i>	35 <i>Dialogue</i>	52 <i>Dialogue</i>
TOC	4 Processes	13 Address Spaces <small>code</small>	26 Concurrency and Threads <small>code</small>	36 I/O Devices	53 <i>Intro Security</i>
1 <i>Dialogue</i>	5 Process API <small>code</small>	14 Memory API	27 Thread API <small>code</small>	37 Hard Disk Drives	54 <i>Authentication</i>
2 Introduction <small>code</small>	6 Direct Execution	15 Address Translation	28 Locks <small>code</small>	38 Redundant Disk Arrays (RAID)	55 <i>Access Control</i>
	7 CPU Scheduling	16 Segmentation	29 Locked Data Structures	39 Files and Directories	56 <i>Cryptography</i>
	8 Multi-level Feedback	17 Free Space Management	30 Condition Variables <small>code</small>	40 File System Implementation	57 <i>Distributed</i>
	9 Lottery Scheduling <small>code</small>	18 Introduction to Paging	31 Semaphores <small>code</small>	41 Fast File System (FFS)	
	10 Multi-CPU Scheduling	19 Translation Lookaside Buffers	32 Concurrency Bugs	42 FSCK and Journaling	<b>Appendices</b>
	11 <i>Summary</i>	20 Advanced Page Tables	33 Event-based Concurrency	43 Log-structured File System (LFS)	<i>Dialogue</i>
		21 Swapping: Mechanisms	34 <i>Summary</i>	44 Flash-based SSDs	<i>Virtual Machines</i>
		22 Swapping: Policies		45 Data Integrity and Protection	<i>Dialogue</i>
		23 Complete VM Systems		46 <i>Summary</i>	<i>Monitors</i>
		24 <i>Summary</i>		47 <i>Dialogue</i>	<i>Dialogue</i>
				48 Distributed Systems	<i>Lab Tutorial</i>
				49 Network File System (NFS)	<i>Systems Labs</i>
				50 Andrew File System (AFS)	<i>xv6 Labs</i>
				51 <i>Summary</i>	

**INSTRUCTORS:** If you are using these free chapters, **please just link to them directly** (instead of making a copy locally); we make little improvements frequently and thus would like to provide the latest to whomever is using it. Also: we have made our own class-preparation notes available to those of you teaching from this book; please drop us a line at [remzi@cs.wisc.edu](mailto:remzi@cs.wisc.edu) if you are interested.

# Homework

- ▣ Some chapter (most) include homework
  - ◆ Homework can be used to solidify your knowledge of the material in each of the chapters
  - ◆ Most homework are based on running little **simulators**, which mimic some aspect of an operating system: For example, a disk scheduling simulator could be useful in understanding how different disk scheduling algorithms work:
    - Most of them provides the **solution**
  - ◆ Some home-works are just short programming exercises, allowing you to explore how real systems work and complement Lab work.
- ▣ Homework are done in **personal-time**

# Labs: C and xv6

- ▣ Refresh C knowledge
- ▣ Use a “toy” kernel to dig into implementation details
  - ◆ It is a clean and beautiful little kernel, and thus a perfect object for our study and usage.
  - ◆ It was developed by OS Eng. In MIT as a port of K&R original Unix R6/PDP11 (6.828 and 6.S081)
  - ◆ Use al real kernel (such as linux) will be certainly overkill

```
cigal xv6-wisc (master) $ cloc *
```

```
145 text files.  
143 unique files.  
15 files ignored.
```

```
http://cloc.sourceforge.net v 1.64 T=1.34 s (99.9 files/s, 8217.7 lines/s)
```

Language	files	blank	comment	code
C	45	946	621	5855
Assembly	9	58	124	1748
C/C++ Header	20	177	138	955
D	57	0	0	154
make	1	40	47	90
Perl	2	11	22	33
SUM:	134	1232	952	8835

# Prerequisites

- ▣ All OS and architecture previous subjects(ugh!)

# Evaluation

- ▣ Full details in:
  - ◆ <http://web.unican.es/estudios/Documents/Guias/2021/es/G677.pdf>
  
- ▣ 20% Parciales
  - ◆ 10% Lab
  - ◆ 10% Teoria
  
- ▣ 80% Examen final
  - ◆ 40% Lab
  - ◆ 40% Teoria



# Schedule

Date	Chapter	Lab	Homework
6-sep.	1 Intro	P0 Lab Intro and review C	
8-sep.	4. The Abstraction: The Process/ 5. Interlude: Process API		Process Intro / Process API
13-sep.	6. Mechanism: Limited Direct Execution		Direct Execution
20-sep.	7. Scheduling: Introduction	P0 Due, P1 System Calls	Scheduler
22-sept	8: Scheduling: The Multi-Level Feedback Queue		MLFQ Scheduling
27-sep.	9: Scheduling: Proportional Share/10. Multiprocessor Scheduling		Lottery Scheduling
29-sep.	13. The Abstraction: Address Space / 14. Memory API	P1 Due, P2 Scheduling	VM API
4-oct.	15. Address Translation		Relocation
6-oct.	16. Segmentation/17. Free-Space Management		Segmentation
11-oct.	18 Pagin Intro.	P2 Due	Free Space
13-oct.	19. Translation Lookaside Buffers	P3 Memory	Paging
18-oct.	20. Paging: Smaller Tables		TLBs
20-oct	21. Swapping: Mechanisms/22. Swaping: Policies		Multi-level Paging/Paging Mechanism
25-oct	26. Concurrency: An Introduction / 27. Interlude: Thread API		Threads (Intro)/Threads (API)
27-oct	28. Locks		Threads (Locks)
3-nov	29. Lock-based Concurrent Data Structures		
8-nov	30. Condition Variables	P3 Due, P4 Threads	Threads (CVs)
???	Mid Term Exam ( Processes & Memory)		
10-nov	31. Semaphore		
15-nov	32. Common Concurrency Problems.		Threads (Bugs)
17-nov	36. I/O Devices		
22-nov	37. Hard Disk Drives		
24-nov	39. File and Directories		Disks
29-nov	40. File system Implementation.	P4 Due, P5 File systems	39. File and Directories
1-dic	41. Fast File System / 42. Crash Consistency: FSCK		FS Implement
13-dic	42. Crash Consistency: Journaling		FFS
15-dic	43. Log-structured File Systems	P5 Due	