

# Sistemas Operativos Avanzados

[vpuente@unican.es](mailto:vpuente@unican.es)

<https://ceunican.github.io/aos/>

<https://www.ce.unican.es/course/soa/>

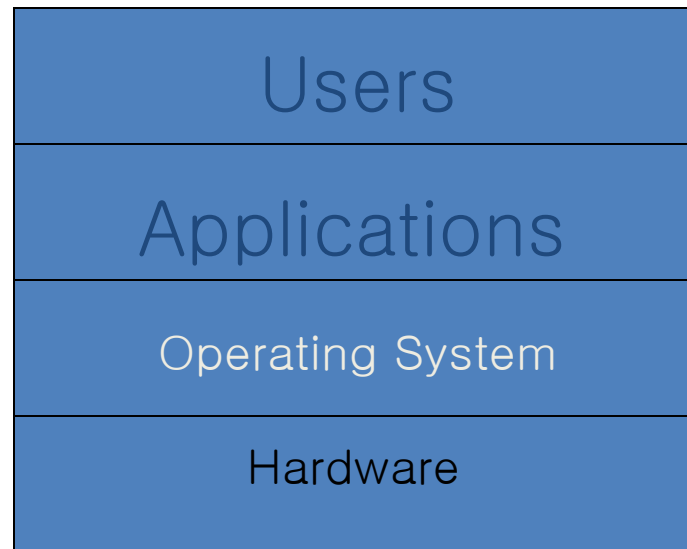


# What is an Operating System?

Operating System (OS):

Software that converts hardware into a useful form for applications

Not easy to define precisely...



# What DOES OS Provide?

- ▣ Role #1: Abstraction - Provide standard library for resources
- ▣ What is a resource?
  - ◆ Anything valuable (e.g., CPU, memory, disk, I/O device)
- ▣ What abstraction does modern OS typically provide for each resource?
  - ◆ CPU:
    - process and/or thread
  - ◆ Memory:
    - address space
  - ◆ Disk:
    - files
- ▣ Advantages of OS providing abstraction?
  - ◆ Allow applications to reuse common facilities
  - ◆ Make different devices look the same
  - ◆ Provide higher-level or more useful functionality
- ▣ Challenges
  - ◆ What are the correct abstractions?
  - ◆ How much of hardware should be exposed?

# What DOES OS Provide?

- ▣ Role #2: Resource management – Share resources well
- ▣ Advantages of OS providing resource management?
  - ◆ Protect applications from one another
  - ◆ Provide efficient access to resources (cost, time, energy)
  - ◆ Provide fair access to resources
- ▣ Challenges
  - ◆ What are the correct mechanisms?
  - ◆ What are the correct policies?

# OS Organization

- ▣ How to cover all the topics relevant to operating systems?

# Three PIECES

- ▣ Virtualization:
  - ◆ Make each application believe it has each resource to itself
  
- ▣ Concurrency:
  - ◆ Events are occurring simultaneously and may interact with one another
  
- ▣ Persistence: Access information permanently
  - ◆ Lifetime of information is longer than lifetime of any one process
  - ◆ Machine may be rebooted, machine may lose power or crash unexpectedly

# Advanced Topics (beyond our reach)

- ▣ Current systems
  - ◆ Multiprocessors
  - ◆ Networked and distributed systems
  - ◆ Virtual machines
  - ◆ Containers
  - ◆ ...
  
- ▣ Many of the pushed by the explosive demand (a.k.a. Massive complexity under constrained cost)
  
- ▣ This is the support of the world: it will keep changing ...
  
- ▣ Some of them covered in SVS (M1679)

# Why study Operating Systems?

- ▣ Build, modify, or administer an operating system
- ▣ Understand system performance
  - ◆ Behavior of OS impacts entire machine
  - ◆ Tune workload performance
  - ◆ Apply knowledge across many layers
    - Computer architecture, programming languages, data structures and algorithms, and performance modeling
- ▣ Fun and challenging to understand large, complex systems
- ▣ **Is the glue that “holds” all the ideas in place**



# Approach

- We will follow the “Operating System: Three Easy Pieces” (OSTEP) style
  - ◆ From the **basic** concepts to state-of-the-art approaches
  - ◆ Eminently **practical** style: all supported by “simulators” and simple coding examples
  - ◆ Assumes some **basic** knowledge in architecture, C, assembler and system administration
  - ◆ More than just a textbook...
  
- Structure
  - ◆ The three parts are split in small *pieces* (~40 in the book)
  - ◆ Each chapter builds on the previous one (can't miss the beat)
  - ◆ Each chapter has attached a “Homework” to reinforce it: from using python simulators to write small pieces of code ( C )
  - ◆ 5 + 1 Labs, to develop on top of **xv6**

# Lecture/Lab structure

- ▣ We mix dynamically both
  - ◆ The real thing is that there is no separation between “theory” and “lab”
  
- ▣ Sessions of:
  - ◆ 1<sup>st</sup> hour: Introduction to the topic
  - ◆ 2<sup>nd</sup> hour: Introduce/develop of Labs
  - ◆ Personal work (out the lab): 6 hours (labs and homework)
  - ◆ 10 hours/week
  - ◆ Strict schedule
  
- ▣ Although the original course/book is designed for 15-week semester (150h work), we will need to drop some details or advanced topics (and half of the labs)

# Material

- ▣ Available in <http://www.ce.unican.es/>
- ▣ All written material will be in “English”
  - ◆ Lecture notes, Homework/Lab guides, etc....
- ▣ *Git* as communication “device”: all material will be delivered via <http://github.com>
  - ◆ An e-mail inviting to join the course project will be sent to unican account
  - ◆ Slides, labs, other reference material is there
  - ◆ It uses “git” to have a “time-track”
    - Lecture notes updates
    - Additional material
- ▣ Use git to allow you and me “track” your personal work
- ▣ Use Github **Discussions** for questions

# Book (ostep.org)

This book is **and will always be free** in PDF form, as seen below. For those of you wishing to **BUY** a copy, please consider the following:



- [Lulu Hardcover \(v1.00\)](#); this may be the best printed form of the book (it really looks pretty good), but it is also the most expensive way to obtain *the black book* of operating systems (a.k.a. *the comet book* or *the asteroid book* according to students). Now just: **\$38.00**
- [Lulu Softcover \(v1.00\)](#); this way is pretty great too, if you like to read printed material but want to save a few bucks. Now just: **\$22.00**
- [Amazon Softcover \(v1.00\)](#); Same book as softcover above, but printed through Amazon CreateSpace. Now just: **\$25.90** (but works with Prime shipping)
- [Downloadable PDF \(v1.00\)](#); this is a nice convenience and adds things like a hyperlinked table of contents, index of terms, lists of hints, tips, systems advice, and a few other things not seen in the free version, all in one massive DRM-free PDF. Once purchased, you will always be able to get the latest version. Just: **\$10.00**
- [Kindle](#); Really, just the PDF and does not include all the bells and whistles common in e-pub books.

**New Partnership:** We have a new partnership with [Educative](#); they offer a way to take an OS course (based on OSTEP) through their platform at a low cost. However, don't worry: the book on this website is and will always remain free.

**Lulu Discount Codes:** These always exist in some form, look around for one?

**Warning:** Some resellers on Amazon buy old versions of the books and claim to sell them as "new" on Amazon (click [here](#) for an example); buy from them at your own risk. In general, buy either directly from Lulu.com or Amazon.com (not a reseller). For Amazon, go to [this page](#) and look for Seller Information to be Amazon.com.

Can't bear to go out in public without OSTEP? How about an [Operating Systems: Three Easy Pieces T-shirt](#) or [laptop sticker](#) or [bathmat](#) or [blanket](#) or [mug](#) or [check out the whole store](#)?

**Donate:** By popular demand, another way to support this site and its contents: **donate!** Click to donate [\\$1](#) - [\\$10](#) - [\\$20](#) - [\\$50](#) - or click [here](#) to donate any amount you want! Your donation helps keep this book going. Think about it: if everyone who came to this website donated just one dollar, we'd have at least three dollars. Thanks!

Another way to help the book out: cite it! Here is the [BiBTeX entry \(seen below\)](#); you can also link to the site of the [best free operating systems book](#) on the market.

## Operating Systems: Three Easy Pieces

Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau  
Arpaci-Dusseau Books  
August, 2018 (Version 1.00)

And now, the free online form of the book, in chapter-by-chapter form (now with chapter numbers!):

Intro	Virtualization		Concurrency	Persistence	Security
Preface	3 <a href="#">Dialogue</a>	12 <a href="#">Dialogue</a>	25 <a href="#">Dialogue</a>	35 <a href="#">Dialogue</a>	52 <a href="#">Dialogue</a>
TOC	4 <a href="#">Processes</a>	13 <a href="#">Address Spaces</a> <small>code</small>	26 <a href="#">Concurrency and Threads</a> <small>code</small>	36 <a href="#">I/O Devices</a>	53 <a href="#">Intro Security</a>
1 <a href="#">Dialogue</a>	5 <a href="#">Process API</a> <small>code</small>	14 <a href="#">Memory API</a>	27 <a href="#">Thread API</a> <small>code</small>	37 <a href="#">Hard Disk Drives</a>	54 <a href="#">Authentication</a>
2 <a href="#">Introduction</a> <small>code</small>	6 <a href="#">Direct Execution</a>	15 <a href="#">Address Translation</a>	28 <a href="#">Locks</a> <small>code</small>	38 <a href="#">Redundant Disk Arrays (RAID)</a>	55 <a href="#">Access Control</a>
	7 <a href="#">CPU Scheduling</a>	16 <a href="#">Segmentation</a>	29 <a href="#">Locked Data Structures</a>	39 <a href="#">Files and Directories</a>	56 <a href="#">Cryptography</a>
	8 <a href="#">Multi-level Feedback</a>	17 <a href="#">Free Space Management</a>	30 <a href="#">Condition Variables</a> <small>code</small>	40 <a href="#">File System Implementation</a>	57 <a href="#">Distributed</a>
	9 <a href="#">Lottery Scheduling</a> <small>code</small>	18 <a href="#">Introduction to Paging</a>	31 <a href="#">Semaphores</a> <small>code</small>	41 <a href="#">Fast File System (FFS)</a>	
	10 <a href="#">Multi-CPU Scheduling</a>	19 <a href="#">Translation Lookaside Buffers</a>	32 <a href="#">Concurrency Bugs</a>	42 <a href="#">FSCK and Journaling</a>	<b>Appendices</b>
	11 <a href="#">Summary</a>	20 <a href="#">Advanced Page Tables</a>	33 <a href="#">Event-based Concurrency</a>	43 <a href="#">Log-structured File System (LFS)</a>	<a href="#">Dialogue</a>
		21 <a href="#">Swapping: Mechanisms</a>	34 <a href="#">Summary</a>	44 <a href="#">Flash-based SSDs</a>	<a href="#">Virtual Machines</a>
		22 <a href="#">Swapping: Policies</a>		45 <a href="#">Data Integrity and Protection</a>	<a href="#">Dialogue</a>
		23 <a href="#">Complete VM Systems</a>		46 <a href="#">Summary</a>	<a href="#">Monitors</a>
		24 <a href="#">Summary</a>		47 <a href="#">Dialogue</a>	<a href="#">Dialogue</a>
				48 <a href="#">Distributed Systems</a>	<a href="#">Lab Tutorial</a>
				49 <a href="#">Network File System (NFS)</a>	<a href="#">Systems Labs</a>
				50 <a href="#">Andrew File System (AFS)</a>	<a href="#">xv6 Labs</a>
				51 <a href="#">Summary</a>	

**INSTRUCTORS:** If you are using these free chapters, **please just link to them directly** (instead of making a copy locally); we make little improvements frequently and thus would like to provide the latest to whomever is using it. Also: we have made our own class-preparation notes available to those of you teaching from this book; please drop us a line at [remzi@cs.wisc.edu](mailto:remzi@cs.wisc.edu) if you are interested.

# Homework

- ▣ Some chapter (most) include homework
  - ◆ Homework will reinforce your understanding of the material covered in each chapter.
  - ◆ Most homework assignments involve running simulators that **mimic** certain aspects of operating systems. For instance, a disk scheduling simulator can help in comprehending the functioning of various disk scheduling algorithms.
    - Most of them provides the **solution**
  - ◆ Some home-works are just short programming exercises, allowing you to explore how real systems work and complement Lab work.
  
- ▣ Homework are done in **personal-time**

# Labs: C and xv6

- ▣ Refresh C knowledge
- ▣ Use a “toy” kernel to dig into implementation details
  - ◆ It is a clean and beautiful little kernel, and thus a perfect object for our study and use.
  - ◆ It was developed by OS Eng. In MIT as a port of K&R original Unix R6/PDP11 (6.828 and 6.S081)
  - ◆ Many possibilities
    - X86, **riscv**, ARM,....
    - C, rust,..

```
cigal xv6-wisc (master)*$ cloc *
145 text files.
143 unique files.
15 files ignored.

http://cloc.sourceforge.net v 1.64 T=1.34 s (99.9 files/s, 8217.7 lines/s)
-----
Language          files      blank      comment      code
-----
C                  45         946         621         5855
Assembly           9           58         124         1748
C/C++ Header       20         177         138          955
D                  57           0           0          154
make               1           40          47           90
Perl               2           11          22           33
-----
SUM:              134        1232         952        8835
-----
```

# Prerequisites

- ▣ All OS and architecture previous subjects(ugh!)

# Evaluation

(<http://web.unican.es/estudios/Documents/Guias/2022/es/G677.pdf>)

T1: Examen Parcial Seguimiento Teoría		Examen escrito	No	Sí	5,00
Calif. mínima	0,00				
Duración	1 hora				
Fecha realización	Semana 9				
Condiciones recuperación	Recuperable realizando el examen final.				
Observaciones	Examen de seguimiento Teoría (Virtualización CPU y Memoria).				
P1: Prácticas de Laboratorio 1		Examen escrito	No	Sí	35,00
Calif. mínima	0,00				
Duración	2,5 horas				
Fecha realización	Semana 9				
Condiciones recuperación	Recuperable realizando el examen final.				
Observaciones	Cuestiones o propuesta de pequeñas modificaciones/extensiones sobre las implementaciones del alumno (prácticas Lab 1, Lab 2, Lab 3 y Lab4)				
T2: Examen Final Teoría		Examen escrito	Sí	Sí	45,00
Calif. mínima	0,00				
Duración	2,5 horas				
Fecha realización	En las fechas indicadas por la Facultad para la realización de exámenes finales				
Condiciones recuperación	Recuperable en la convocatoria extraordinaria				
Observaciones	Preguntas que evaluarán globalmente el grado de comprensión de la materia de la asignatura.				
P2: Prácticas de Laboratorio 2		Examen escrito	Sí	Sí	15,00
Calif. mínima	0,00				
Duración	1,5 horas				
Fecha realización	En las fechas indicadas por la Facultad para la realización de exámenes finales				
Condiciones recuperación	Recuperable en la convocatoria extraordinaria				
Observaciones	Cuestiones o propuesta de pequeñas modificaciones/extensiones sobre las implementaciones del alumno (prácticas Lab5,y Lab6)				

Para poder superar la asignatura, las notas medias de la parte práctica (i.e,  $P1*0.7+P2*0.3$ ) y la parte teórica (i.e.,  $T1*0.1+T2*0.9$ ) deberán ser superior a 3.0.



# Schedule (tentative)

2023/2024	Chapter	Lab	Homework
4-sep.	1 Intro	P0 Lab Intro and review C	
6-sep.	4. The Abstraction: The Process/ 5. Interlude: Process API		Process Intro / Process API
11-sep.	6. Mechanism: Limited Direct Execution		Direct Execution
13-sep.	7. Scheduling: Introduction	P0 Due, P1 System Calls	Scheduler
18-sept	8: Scheduling: The Multi-Level Feedback Queue		MLFQ Scheduling
20-sep.	9: Scheduling: Proportional Share/10. Multiprocessor Scheduling		Lottery Scheduling
25-sep.	13. The Abstraction: Address Space / 14. Memory API	P1 Due, P2 Scheduling	VM API
27-sep.	15. Address Translation, 16. Segmentation		Relocation
2-oct.	17. Free-Space Management		Segmentation
4-oct.	18 Pagin Intro.	P2 Due	Free Space
9-oct.	19. Translation Lookaside Buffers	P3 Memory	Paging
11-oct	20. Paging: Smaller Tables		TLBs
16-oct	21. Swapping: Mechanisms/22. Swaping: Policies		Multi-level Paging/Paging Mechanism
18-oct	26. Concurrency: An Introduction / 27. Interlude: Thread API		Threads (Intro)/Threads (API)
23-oct	28. Locks		Threads (Locks)
25-oct	29. Lock-based Concurrent Data Structures		
30-oct	30. Condition Variables	P3 Due, P4 Threads	Threads (CVs)
	Mid Term Exam ( Processes & Memory LAB & TEO)		
6-nov	32. Common Concurrency Problems.		
8-nov	36. I/O Devices		Threads (Bugs)
13-nov	37. Hard Disk Drives		
15-nov	39. File and Directories		
20-nov	40. File system Implementation.		Disks
22-nov	41. Fast File System / 42. Crash Consistency: FSCK	P4 Due, P5 File systems	39. File and Directories
27-nov	42. Crash Consistency: Journaling		FS Implement
29-nov	43. Log-structured File Systems		FFS
4-dic	44. SSD		
11-dic	41.RAID		
13-nov	45. Data Integrity And Protection	P5 Due	